

Introduction to Machine Learning

Module 1: Foundations

Tobias Rebholz

University of Tübingen

Fall 2024, SMiP Workshop

EBERHARD KARLS
UNIVERSITÄT
TÜBINGEN



Orga

Learning Target

- Develop a basic understanding of ML and its value to applied psychological research
 - Emphasis is on conceptual understanding, rather than (technical) details!
- Gain practical experience in applying specific ML techniques to solve real-world decision problems
 - Extended skills to analyze and interpret complex data sets, such as unstructured text data
- Acquire the competence to critically reflect on the results of ML methods and to evaluate their implications for theory building and psychological research practice
- Prerequisites:
 - Basic understanding of statistics (e.g., linear regression)
 - Familiarity with R

Literature

1. James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). *An introduction to statistical learning: With applications in R* (2nd ed.). Springer US. <https://doi.org/10.1007/978-1-0716-1418-1>
 - More technical, less related to behavioral science, but more standard
 - Available for free at: <https://www.statlearning.com/>
 - Now also available for Python
2. Jacobucci, R., Grimm, K. J., & Zhang, Z. (2023). *Machine Learning for social and behavioral research* (2nd ed.). The Guilford Press.
 - Less technical, more related to behavioral science, but less standard
 - Should be freely available in Ba-Wü: E.g., via proquest
3. Hastie, T. J., Tibshirani, R., & Friedman, J. H. (2009). *The elements of statistical learning: Data mining, inference, and prediction* (2nd ed.). Springer.
 - Extremely technical, but basically the bible for ML

Structure

The workshop is divided into two parts:

- Introduction to topics/methods
 - Mainly lectured by me!
- Solving simple programming assignments
 - Mainly hands-on by you!

Material

- Slides, programming assignments etc. will be provided via my personal website:
<https://tobiasrebholz.github.io/teaching/>
- Password: ...

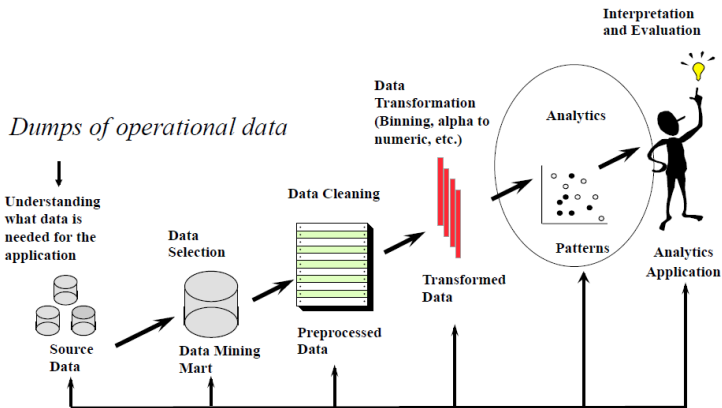
Tentative Schedule

- Zeitplan gestalten

What is Data-Driven Decision-Making?

- Data Science aims to extract and represent knowledge from complex data
 - Machine Learning refers to a vast set of mathematical tools and models which can help to make sense of data
- Techniques from diverse fields, such as:
 - Statistics,
 - Data Mining
 - Visualization
 - ...
- Expertise from different disciplines, such as:
 - Statistics
 - Computer Science
 - Behavioral Science
 - Neuroscience
 - ...

Data Science enables Data-Driven Decision-Making



(https://ebrary.net/168827/computer_science/twitter_s_data_analysis_using_rstudio)

Examples: Medical Diagnosis

AI Now Diagnoses Disease Better Than Your Doctor, Study Finds

Peer-reviewed study says you'll soon consult Dr. Bot for a second opinion

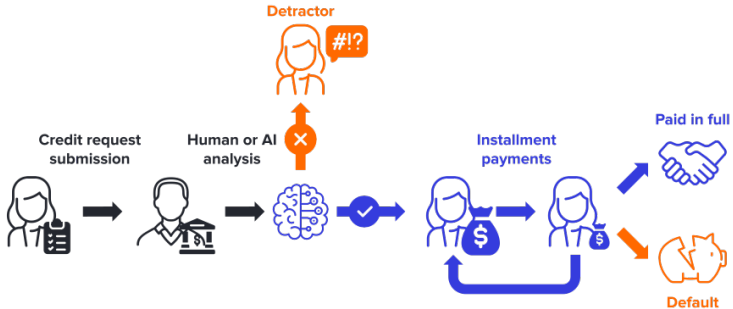


David Leibowitz · Follow

Published in Towards Data Science · 7 min read · Sep 29, 2020


- “In the peer-reviewed study, authored by researchers from Babylon Health and University College London, the new model scored higher than 72% of general practitioner doctors when tasked with diagnosing written test cases of realistic illnesses.” (<https://towardsdatascience.com/ai-diagnoses-disease-better-than-your-doctor-study-finds-a5cc0ffbf32>)

Examples: Credit Scoring in FinTech Industry



<https://nilg.ai/202107/insights-in-ai-applied-to-credit-scoring/>

Examples: Sports



<https://en.wikipedia.org/wiki/Moneyball>

A Brief History of ML

- Early 19th century: Method of least squares was developed
 - Precursor to linear regression
- 1940s: Introduction of logistic regression for predicting qualitative values
- Early 1970s: Coining of the term “generalized linear models”, encompassing linear and logistic regression
- 1980s: Improvement in computing technology allows for nonlinear methods
- Mid-1980s: Development of classification and regression trees
- 1980s: Early neural networks (e.g., perceptron)
- 1990s: Emergence of support vector machines
- 2001: Random Forests (i.e., ensemble methods)
- 2010s: Deep Learning (i.e., advanced neural networks)
- 2017: Transformer models (e.g., large language models)
 - Enabled high-performance chatbots, such as OpenAI’s ChatGPT

Analytics

- Supervised Learning: Predicting/estimating an output based on one or more inputs
- Unsupervised Learning: Learning relationships/structure from inputs only
 - I.e., there is no supervising output
- Visualization: Visual representation of predictions, relationships, but also raw data

Data Types

- Continuous data
 - E.g., age, income, ...
 - Can be discretized (e.g., age categories)
- Binary data
 - E.g., yes vs. no response
 - Usually coded as a 0-1 dummy variable
- Categorical data
 - E.g., gender, group membership, ...
 - Sometimes converted into dummies: One dummy variable per category
- Ordinal data
 - E.g., highest educational degree
 - Must be treated with caution, as they are neither continuous nor categorical

Complex Data: Multivariate

- Data on demographics and Big Five personality trait scores:

head(dat, 15)										
##	CASE	gender	education	age	agree	conscientious	extra	neuro	open	
## 652	63116	1	3	40.55747	5.8	3.40	3.6	1.5	3.80	
## 999	63967	1	4	35.37340	4.6	3.60	4.0	1.0	3.60	
## 991	63955	2	4	21.55923	3.0	3.20	4.0	3.8	4.40	
## 392	62547	2	1	22.80091	4.8	5.20	4.4	2.0	4.80	
## 788	63493	2	2	23.07484	5.2	3.40	4.4	2.6	4.60	
## 330	62419	2	3	30.08120	5.4	4.00	4.4	4.0	3.80	
## 2231	66716	2	5	37.61331	3.8	4.80	4.4	4.2	5.00	
## 1128	64275	1	4	37.76547	5.0	5.00	4.6	3.6	4.60	
## 1061	64101	1	3	17.58908	4.6	4.00	4.6	1.8	3.60	
## 1474	65080	2	4	22.49090	2.8	3.80	4.0	3.0	3.00	
## 1949	66088	1	2	58.15496	5.0	4.75	5.0	4.8	5.00	
## 1005	63983	1	3	23.31878	3.6	4.60	3.2	2.6	3.40	
## 261	62266	1	3	30.89218	3.0	3.40	3.6	4.0	4.00	
## 2024	66264	2	3	24.58253	4.8	3.40	4.2	2.0	4.75	
## 2441	67257	2	5	18.87679	4.8	4.40	4.0	4.6	4.00	

Complex Data: Spatial


- Traffic data:



<https://unece.org/traffic-census-map>

Complex Data: Time-stamped

- Stock market data:



<https://www.r-bloggers.com/2020/05/forecasting-the-next-decade-in-the-stock-market-using-time-series-models/>

Complex Data: Relational


- Network data:

Episode III: Revenge of the Sith




Open network

Episode IV: A New Hope




Open network

Episode V: The Empire Strikes Back



Open network


Episode VI: Return of the Jedi



<https://www.martingrandjean.ch/star-wars-data-visualization/>

Complex Data: Unstructured

- Image data:



<https://github.com/zandreika/letters-recognition>

- Further unstructured data: Videos, audios, text, health records, ...

Statistics of Data

- Mean: $\mu = \frac{1}{N} \sum_{i=1}^N x_i$
- Variance: $\sigma^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2$
- Standard deviation: $\sigma = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2}$
- Minimum: $\min = \text{minimize}_{i \in N} \{x_i\}$
- Maximum: $\max = \text{maximize}_{i \in N} \{x_i\}$
- Bringing variables to the same range (very useful for multivariate data):
 - Standarizing:

$$\frac{x - \mu}{\sigma} \quad (1)$$

- Normalizing:

$$\frac{x - \min}{\max - \min} \quad (2)$$

Summary

- Data Science builds mathematical models to extract and represent knowledge
- Data-Driven Decision-Making can be relevant to any setting in:
 - Psychology
 - Economics
 - Medicine
 - Physics
 - ...
- Prior to building the models, data often needs some preprocessing

Linear Regression

Research Goal

- **Description:** Research with the aim of describing relationships or distributions

VS.

- **Explanation:** Research with the aim of understanding the underlying mechanisms


VS.

- **Prediction:** Research with the aim of maximally explaining the variability in an outcome
 - ML, and therefore this course, is mainly devoted to predictive analytics

Simple Linear Regression

```
mdl <- lm(agree ~ age, data = dat)

mdl %>%
  ggplot(aes(y = agree, x = age)) +
  geom_point() +
  geom_smooth(method = "lm") +
  geom_segment(aes(xend = age, yend = .fitted))
## `geom_smooth()` using formula = 'y ~ x'
```



Simple Linear Regression in base R

```
summary mdl)
##
## Call:
## lm(formula = agree ~ age, data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9057 -0.5915  0.1043  0.5542  1.5251
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.373290   0.223798   15.073 < 2e-16 ***
## age          0.028270   0.007092    3.986 0.000129 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7274 on 98 degrees of freedom
## Multiple R-squared:  0.1395, Adjusted R-squared:  0.1307
## F-statistic: 15.89 on 1 and 98 DF,  p-value: 0.0001292
```

- Fitted model: $\hat{agree} = \hat{\beta}_0 + \hat{\beta}_1 * age = 3.37 + 0.03 * age$
 - Prediction for a new participant with $age = 30$:
 $\hat{agree} = 3.37 + 0.03 * 30 = 4.22$

Prediction “Out-of-Sample”

1. Use a subset of the sample to fit the model:

```
N <- nrow(dat)
train <- sample(1:N, N*0.9)

df_subset <- dat[train,]

mdl_subset <- lm(agree ~ age, data = df_subset)
summary(mdl_subset)
##
## Call:
## lm(formula = agree ~ age, data = df_subset)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9365 -0.6068  0.1003  0.5975  1.5335
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.440593   0.246566  13.954 < 2e-16 ***
## age          0.026329   0.007975   3.301  0.00139 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7449 on 88 degrees of freedom
## Multiple R-squared:  0.1102, Adjusted R-squared:  0.1001
```

Prediction “Out-of-Sample”

2. Test the model on the remaining, held-out data

```
df_rest <- dat[-train,]

df_rest$agree_predicted <- predict(mdl_subset, df_rest)

df_rest %>% select(agree, agree_predicted)
##      agree agree_predicted
## 1949    5.0      4.971744
## 274     2.8      3.939006
## 2378    4.8      4.560227
## 1672    5.0      4.204977
## 1000    4.2      3.970284
## 737     4.8      4.644516
## 1754    4.8      4.621869
## 2550    4.0      4.106741
## 816     3.2      4.079394
## 744     4.0      4.437448
```

- We will expand on this idea over the weeks

Simple Linear Regression in mlr3

1. Define a (prediction) task, which is mlr3's way to store the raw data along with some meta-information for modeling
2. Specify which ML model to apply later for prediction
 - mlr3 does not implement its own ML models but links to available implementations in other R packages (e.g., "regr.lm" links to the ordinary `lm()` function in the stats package)
3. Train the linear regression model
 - In mlr3, objects have "abilities" (or "methods") that can be applied with the following `$`-syntax (here, the train method of the learner object is used to train the learner from step 2 on the task specified in step 1)

```
library(mlr3verse)
tsk = as_task_regr(agree ~ age, data = dat) #1.
mdl = lrn("regr.lm") #2.
mdl$train(tsk) #3.
```

Simple Linear Regression in mlr3

- The estimated model output is exactly the same as with using base R
 - Not surprising, as `lrn("regr.lm")` is simply applying the `lm()` function to estimate the model, which is exactly equivalent to what we did a couple of slides ago

```
summary(mdl$model)
##
## Call:
## stats::lm(formula = task$formula(), data = task$data())
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9057 -0.5915  0.1043  0.5542  1.5251
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.373290   0.223798  15.073  < 2e-16 ***
## age          0.028270   0.007092   3.986  0.000129 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7274 on 98 degrees of freedom
## Multiple R-squared:  0.1395, Adjusted R-squared:  0.1307
## F-statistic: 15.89 on 1 and 98 DF, p-value: 0.0001292
```


Excuse: Why mlr3?

- If the end result (i.e., fitted model) is exactly the same as with base R, why use mlr3 at all?
- In mlr3, the model fitting procedure (i.e., steps 1-3) is exactly the same **no matter what ML method we're using** to model the data
 - It merely provides a **unified framework/syntax** (cf. tidyverse) to fit different ML models
 - The ML models themselves, however, stem from other, established R packages (e.g., “regr.lm” links to the ordinary `lm()` function in the stats package), which are also described in detail in, e.g., James et al. (2021)
 - Allows us to **focus more on the concepts** and less on the specific programming in R
- mlr3 is **the state-of-the-art framework for ML in R** (cf. scikit-learn in Python)

Excuse: Nonlinear Regression


- Using linear modeling, we can even build nonlinear models
 - E.g., by including polynomial transformations of features (i.e., quadratic age): $\hat{agree} = \hat{\beta}_0 + \hat{\beta}_1 * age + \hat{\beta}_2 * age^2$

```
ggplot(dat, aes(y = agree, x = age)) +  
  geom_point() +  
  stat_smooth(method = "lm", formula = y ~ x + I(x^2))
```



Bias-Variance Trade-Off


- Bias-variance trade-off: Good out-of-sample performance requires low variance as well as low squared bias
 - Why “trade-off”? – “Inflexible” model with low variance but high bias vs. “flexible” model with low bias but high variance
 - The challenge lies in finding a model for which both the variance and the (squared) bias are low



(Pargent et al., 2023, Figure 3a)

Bias-Variance Trade-Off

- Darts metaphor:
 - Bullseye = True value (i.e., expected target value for a given combination of feature values)
 - Each cross is the prediction made by a concrete ML model (trained on a randomly drawn training set, all from the same population and with the same sample size)




(Pargent et al., 2023, ESM, Figure 2)

- Optimal model: Low bias and low variance (bottom left)
 - Noise = Irreducible error of the true model: Reason why predictions (across different samples) are not similar, even when hitting the bullseye

Multiple Linear Regression

- E.g., vector of features of size 2:



(James et al., 2021, Figure 3.4)

Multiple Linear Regression in base R

```
mdl <- lm(agree ~ age + gender, data = dat)
summary(mdl)
##
## Call:
## lm(formula = agree ~ age + gender, data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.86767 -0.49583  0.07832  0.53912  1.45493
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.040273   0.351015   8.661 1.04e-13 ***
## age          0.028927   0.007093   4.078 9.31e-05 ***
## gender       0.188803   0.153585   1.229  0.222
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7255 on 97 degrees of freedom
## Multiple R-squared:  0.1527, Adjusted R-squared:  0.1353
## F-statistic: 8.743 on 2 and 97 DF,  p-value: 0.0003229
```

Multiple Linear Regression in mlr3

```
tsk = as_task_regr(agree ~ age + gender, data = dat)
mdl = lrn("regr.lm")
mdl$train(tsk)
summary(mdl$model)
##
## Call:
## stats::lm(formula = task$formula(), data = task$data())
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.86767 -0.49583  0.07832  0.53912  1.45493
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.040273   0.351015   8.661 1.04e-13 ***
## age          0.028927   0.007093   4.078 9.31e-05 ***
## gender       0.188803   0.153585   1.229  0.222
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7255 on 97 degrees of freedom
## Multiple R-squared:  0.1527, Adjusted R-squared:  0.1353
## F-statistic: 8.743 on 2 and 97 DF,  p-value: 0.0003229
```

Hands-on Practical Tutorial

- Now it's your turn:
 - Go to ILIAS
 - Navigate to the "Tutorials" folder
 - Download the "Module 1" folder
 - Work on the tutorial "module1-linear_regression"

Logistic Regression

Logistic Regression

- Auxiliary target: Probability of membership in class 1, p
 - Counter probability $1 - p$: Membership in class 0
 - **Continuous, but bounded** target
- Auxiliary goal: Predict the auxiliary target (i.e., probability) for new instances for whom we know the vector of features but not the value of the target:

$$X_{new} \rightarrow \hat{p}_{new} \in (0, 1) \quad (7)$$

- Predicted class:

$$\hat{y}_{new} = \begin{cases} 1 & \text{if } \hat{p}_{new} > 0.5 \\ 0 & \text{else} \end{cases} \quad (8)$$

Logistic Regression

- Fitting a logistic function to the probability of class 1:

$$p = \frac{e^{\beta_0 + \beta_1 x_1 + \dots + \beta_p x_p}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots + \beta_p x_p}} \quad (9)$$

... is equivalent to fitting a linear function to the logarithm of the odds

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p \quad (10)$$

- In general, “odds” is defined as the probability of an event occurring relative to the probability of the event not occurring
 - Here, the odds are the probability for membership in class 1 relative to the probability for membership in class 0
- As $\beta_0, \beta_1, \dots, \beta_p$ are unknown, we have to estimate them from the data
 - E.g., by maximizing the log likelihood function $\Rightarrow \hat{\beta}_p$ is called the “maximum likelihood estimate” (MLE)

Logistic Regression in mlr3

- Data preparation:

```
dat$agree_high <- ifelse(dat$agree > 4, 1, 0)
dat %>% select(agree, agree_high) %>% tail(., 10)
```

##	agree	agree_high
## 442	4.8	1
## 2270	5.0	1
## 1133	5.2	1
## 1210	4.8	1
## 1912	4.6	1
## 744	4.0	0
## 1485	5.0	1
## 9	4.0	0
## 866	3.6	0
## 1591	4.4	1

Logistic Regression in mlr3

- Model fitting:

```
tsk = as_task_classif(agree_high ~ age, data = dat, positive = "1")
mdl = lrn("classif.log_reg")
mdl$train(tsk)
summary(mdl$model)
##
## Call:
## stats::glm(formula = task$formula(), family = "binomial", data = data,
##      model = FALSE)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.83118      0.73681  -2.485  0.01294 *
## age          0.07940      0.02561   3.100  0.00193 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 133.75  on 99  degrees of freedom
## Residual deviance: 121.83  on 98  degrees of freedom
## AIC: 125.83
##
## Number of Fisher Scoring iterations: 4
```

Logistic Regression in mlr3

- Fitted model: $\hat{p}_{agree>4} = \frac{e^{\hat{\beta}_0 + \hat{\beta}_1 * age}}{1 + e^{\hat{\beta}_0 + \hat{\beta}_1 * age}} = \frac{e^{-1.83 + 0.08 * age}}{1 + e^{-1.83 + 0.08 * age}}$
- Prediction for a new participant with $age = 30$:

$$\hat{p}_{agree>4} = \frac{e^{-1.83 + 0.08 * 30}}{1 + e^{-1.83 + 0.08 * 30}} = 0.63$$

```
summary mdl$model$coefficients
##           Estimate Std. Error  z value    Pr(>|z|)
## (Intercept) -1.83118304 0.73680914 -2.485288 0.012944661
## age          0.07940429 0.02561037  3.100474 0.001932111
```


- Odds ratios:

```
exp(coefficients(mdl$model))
## (Intercept)          age
##    0.1602239    1.0826419
```

Logistic Regression in mlr3

- We can plot the estimated probability of class 1 membership as a function of age:

```
ggplot(dat, aes(y = agree_high, x = age)) +  
  geom_point() +  
  geom_smooth(method = "glm", method.args = list(family = binomial(link = "logit")))  
## `geom_smooth()` using formula = 'y ~ x'
```



Classification Performance

- Main goal: Correct classification
 - E.g., minimizing the mean misclassification error:

$$MMCE = \frac{1}{N} \sum_{i=1}^N I\{y_i \neq \hat{y}_i\} \quad (11)$$

- Where $I\{\cdot\}$ is the indicator function taking the value 1 if the condition in the parentheses is true and 0 otherwise
- This is equivalent to maximizing classification accuracy:

$$Acc = 1 - MMCE = \frac{N_{0,0} + N_{1,1}}{N} \quad (12)$$

- Confusion matrix:

	$\hat{y} = 0$	$\hat{y} = 1$	
$y = 0$	$N_{0,0}$ (TN)	$N_{0,1}$ (FP)	$N_{0,\cdot}$
$y = 1$	$N_{1,0}$ (FN)	$N_{1,1}$ (TP)	$N_{1,\cdot}$
	$N_{\cdot,0}$	$N_{\cdot,1}$	N

Classification Performance

- Other important metrics:
 - Sensitivity (or true positive rate, recall):


$$Sens = \frac{TP}{TP + FN} \quad (13)$$

- Specificity (or true negative rate):

$$Spec = \frac{TN}{TN + FP} \quad (14)$$



Classification Performance

- Predicted class: $\hat{y} = \begin{cases} 1 & \text{if } \hat{p} > 0.5 \\ 0 & \text{else} \end{cases}$
 - Other thresholds than 0.5 can lead to better performance
- Area under the (receiver operating) curve (AUC):
 - The probability that an observation randomly drawn from class 1 has a higher predicted probability to belong to class 1 than an observation randomly drawn from class 0
 - The ROC traces out *Sens* and *Spec* for varying classification thresholds



Excuse: Logistic vs. Linear Regression

- Linear regression: Some estimated probabilities are negative
- Logistic regression: All estimated probabilities lie between 0 and 1 (i.e., are well-defined)

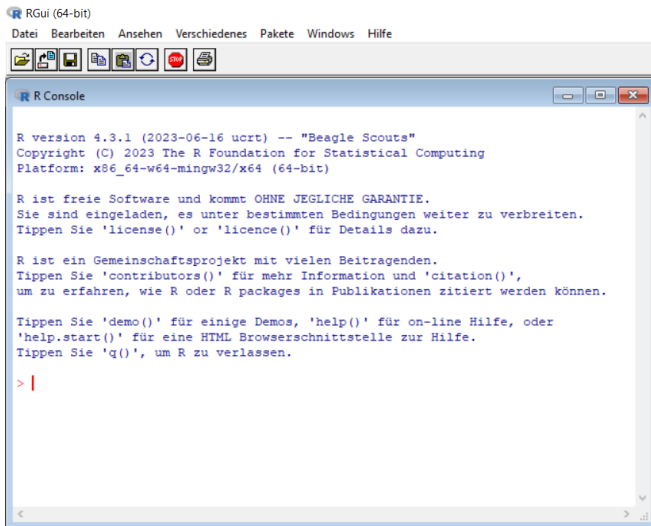


(James et al., 2021, Figure 4.2)

Summary

Appendix

R Console



```
RGui (64-bit)
Datei Bearbeiten Ansehen Verschiedenes Pakete Windows Hilfe

R Console

R version 4.3.1 (2023-06-16 ucrt) -- "Beagle Scouts"
Copyright (C) 2023 The R Foundation for Statistical Computing
Platform: x86_64-w64-mingw32/x64 (64-bit)

R ist freie Software und kommt OHNE JEGliche GARANTIE.
Sie sind eingeladen, es unter bestimmten Bedingungen weiter zu verbreiten.
Tippen Sie 'license()' or 'licence()' für Details dazu.

R ist ein Gemeinschaftsprojekt mit vielen Beitragenden.
Tippen Sie 'contributors()' für mehr Information und 'citation()',
um zu erfahren, wie R oder R packages in Publikationen zitiert werden können.

Tippen Sie 'demo()' für einige Demos, 'help()' für on-line Hilfe, oder
'help.start()' für eine HTML Browserschnittstelle zur Hilfe.
Tippen Sie 'q()', um R zu verlassen.

> |
```

Installation

- Go to <http://www.r-project.org>
 - Click on the download link on the main page
 - Choose your preferred CRAN mirror
 - Navigate to “Download and Install R”
 - Choose your platform
-
- Important: We will be working with R Version 4.3.3
 - Same version makes troubleshooting much easier
 - You can check your installed version by executing `R.Version()` in your console

RStudio Editor

The screenshot displays the RStudio Editor interface with four main panes:

- Source:** Contains R code for creating a ggplot2 plot. The code is:


```
1 library(ggplot2)
2 mpg_plot <- ggplot(mpg, aes(x = displ, y = hwy)) +
3   geom_point(aes(colour = class))
4
5 mpg_plot
6
```
- Environment:** Shows the Global Environment with a table of objects:

Name	Type	Len	Size	Value
mpg_plot	gg	9	29.1	List of 9
- Console:** Shows the execution of the code in the Source pane:


```
R 4.2.0 ~ /rstudio-user-guide/
> library(ggplot2)
> mpg_plot <- ggplot(mpg, aes(x = displ, y = hwy)) +
+   geom_point(aes(colour = class))
>
> mpg_plot
>
```
- Plots:** Displays a scatter plot of highway mileage (hwy) versus engine displacement (displ), colored by vehicle class. The legend indicates the following classes: 2seater, compact, midsize, minivan, pickup, subcompact, and suv.

Coding in R

- Command lines start with ">"
- Comment lines start with "#"
- Use "<-" or "=" to assign a value to a variable
- Use "==", "<", ">", and "!=" to check logical conditions

Variables, Vectors, Matrices, ...

- Variables: Define a variable, x , assign it the value 1, print x , add 5 to x , print x again

```
x <- 1
print(x)
## [1] 1

x <- x + 5
print(x)
## [1] 6
```

- Vectors: Define two vectors, $u = (1, 4, 10, -1)$ and $v = (10, -4, 3, 0)$, print $u + v$

```
u <- c(1, 4, 10, -1)
v <- c(10, -4, 3, 0)

z = u + v
print(z)
## [1] 11 0 13 -1
```


- Statistical functions:

```
u
## [1] 1 4 10 -1

mean(u)
## [1] 3.5

sd(u)
## [1] 4.795832

median(u)
## [1] 2.5

summary(u)
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      -1.0    0.5      2.5      3.5    5.5     10.0
```

Dataset Handling

- Reading from/writing to a file:

```
#reading data from drive
dat <- read.csv('subfiles/data/bfi.csv', header = TRUE)
head(dat)
##      CASE gender education      age agree conscientious extra neuro open
## 1 63116      1         3 40.55747   5.8          3.4    3.6    1.5    3.8
## 2 63967      1         4 35.37340   4.6          3.6    4.0    1.0    3.6
## 3 63955      2         4 21.55923   3.0          3.2    4.0    3.8    4.4
## 4 62547      2         1 22.80091   4.8          5.2    4.4    2.0    4.8
## 5 63493      2         2 23.07484   5.2          3.4    4.4    2.6    4.6
## 6 62419      2         3 30.08120   5.4          4.0    4.4    4.0    3.8

# writing data to drive
write.csv(dat, file = 'subfiles/data/bfi.csv', row.names = FALSE)
```


- Well-known data repositories for research purposes:
 - <http://archive.ics.uci.edu/ml/>
 - <https://www.kaggle.com/datasets>

Plotting

- Histogram:

```
hist(dat$agree)
```

Histogram of dat\$agree



Packages

- Some packages are included by default, such as stats and graphics
 - Use `installed.packages()` to find out which ones
- Other packages can be installed using `install.packages("name")`
 - Packages provide extended functionality for R
 - Most important package for data handling: `tidyverse`
 - Most important package for ML: `mlr3verse`
 - Ideally, install packages with `dependencies = TRUE` to also include any packages that might be used in the package name you are installing
- You have to load new packages with `library("name")` every time you start a new session
 - Alternatively, you can access individual functions from installed packages without loading the entire package (cf. Python) with `name::function()`

Packages

- Caution: Different packages can have the same names for different functions
 - The last loaded package overwrites any functions with the same name in previously loaded packages
- Recommendations:
 - Load all packages needed for the analysis in one section at the beginning of the script
 - Usually it is best to load `tidyverse` and `mlr3verse` last
 - Pay attention to warnings when loading packages
 - If certain functions are overwritten, you can still access them with `name::function()`
 - Always start a new R session when you start your analysis to avoid unexpected behavior from packages loaded from previous sessions

More to come

- You can access the internal help with `?command`
- Google (and recently also GPT) are your best friends for solving programming issues
- With the practical tutorials, we will get more familiar with R